Our Backyard Is Not Theirs: On LLMs and the Outsourcing of Virtue

Jonathan A. Karr, Jr.
University of Notre Dame
jkarr@nd.edu

Nitesh V. Chawla
University of Notre Dame
nchawla@nd.edu

Abstract

Large Language Models (LLMs) are becoming embedded in how we work, learn, and navigate our moral and spiritual lives. While they offer utility in retrieving information and solving technical problems, users often accept their outputs without thinking critically. This has led users to treat statistically generated responses as authoritative guidance. Drawing on Catholic Social Teaching, LLMs should not be treated as an AI companion that solves our moral questions because we should not outsource virtue to them. LLMs should be properly understood as tools, such as a shovel. They are instruments of human dominion and can assist us with stewardship over creation, capable of serving human flourishing or undermining it. To treat them as moral companions is to forgo the very exercise through which conscience is cultivated. At the end of the day, we are the ones asking the questions, and it is our duty to use them responsibly.

Keywords: large language models, virtue ethics, Catholic Social Teaching, moral agency, human flourishing

1. Introduction

When a person encounters a difficult moral question, they have traditionally turned to trusted sources of guidance: family, friends, mentors, and spiritual directors. However, now it has become easy to open a chat window to get the answer in seconds. Large Language Models (LLMs) have become part of daily life, are here to stay, and are a sign of the times. Yet, they pose a critical problem. They respond in a way that pleases the user (Zhao et al., 2025), answer with data that is trained on the masses, and their speed has outpaced our philosophical reflection. We can become tempted to use them when we should not, since they sit in our pockets, answer our questions, and serve as the first point of contact for a growing number of users (Xu et al., 2024). We are asking questions that are not merely technical but moral, spiritual, and deeply human. LLMs have become, in a meaningful sense, our backyard: the intimate space where we bring what troubles us. We argue that virtue should not be outsourced to them. The concern is not primarily that LLMs give wrong answers. Nor is it simply that users fail to think critically. The deeper problem is one that virtue ethics uniquely identifies: virtue is not a product to be retrieved but a capacity formed through repeated moral action, genuine encounter with others, and practical wisdom that only develops through exercise. Every time a person delegates a moral question to an LLM rather than wrestling with it, they forgo an opportunity to cultivate the very habit they need. The damage is not in the output. It is in the habituation which prohibits us from forming our conscience, and this is essential to human flourishing (United States Conference of Catholic Bishops, 2006, p. 314).

We call this dynamic **moral outsourcing** and argue it represents a structural threat to human moral agency. Drawing on our two prior empirical studies, we see that LLMs equivocate and erode Catholic theological and moral truth (Karr Jr et al., 2026), and see how the NIMBYist bias against people experiencing homelessness, encoded in social media platforms on which LLMs are trained (Karr Jr et al., 2026), is reproduced in LLM outputs. Therefore, this is not a hypothetical risk. It is already happening. LLMs do not merely fail

to provide virtue. They risk displacing it.

We borrow from the NIMBY phenomenon, "Not In My Backyard", in which people support a social good in principle while refusing its demands of subsidiarity: the local, inconvenient, embodied encounter that authentic commitment requires. We argue that LLM users exhibit an analogous structure: they retain a nominal commitment to moral life while delegating its difficult exercise to a machine. We must remember that the backyard is an intimate space of moral formation that is ours. Therefore, our backyard shall not be outsourced to LLMs.

To make this argument, we draw on Catholic Social Teaching (CST) alongside the Aristotelian and Thomistic virtue ethics tradition that grounds it. The question virtue ethics asks is not whether the tool produces good outcomes, but what kind of moral agent you are becoming in your use of it. That is the core question we examine in our paper.

So what exactly does it mean that "Our Backyard is not Theirs"? Our backyard is our backyard as humans, and their backyard is the backyard of LLMs. We form this in three steps.

1. Our Backyard: LLMs are becoming integrated into our lives through daily use. It is becoming our neighbor and who we talk to. It is our playground where we ask questions and seek answers.

2. Not in My Backyard: The NIMBY phenomenon is common on social media platforms, and it is not just humans who are amplifying this data. LLMs are trained on social media platforms such as Reddit, which have been shown to amplify these biases (Karr Jr et al., 2026).

3. Our Backyard is Not Theirs: LLMs are trained on data from the internet; yet our conscience as humans is formed through personal experience and interactions with others. This fundamental divide is important to remember, and makes our backyard not theirs.

We outline the phenomenon of moral outsourcing and introduce the NIMBY parallel in Section 2. In Section 3, we draw on empirical findings on equivocation and erosion in

LLM theological discourse, including the limits of purpose-built Catholic LLM. Next, we examine NIMBYist bias in LLM training data and outputs in Section 4. Then, in Section 5, we ground our findings the virtue ethics and the CST framework. After, we defend the instrumentalist account of LLMs and address counterarguments. In Section 7, we explore what responsible use looks like. Finally, we conclude by returning to the backyard and to the irreducible human duty to tend it.

2. The Backyard Problem: LLMs as Moral Companions

LLMs are not merely search engines that get information. They are increasingly positioned by their developers, their users, and the cultural moment as companions. Users disclose anxieties to them, seek relationship advice, process grief, and ask what they should do when they do not know what to do. The scale of this shift among young people is striking. According to a 2025 Common Sense Media study, 72% of teens aged 13 to 17 have used artificial intelligence (AI) companions, 52% are regular users, and nearly one-third find AI conversations as satisfying or more satisfying than human conversations (Robb and Mann, 2025). A separate Pew Research survey found that approximately three in ten teens report using AI chatbots every day, with 16% doing so several times a day (Faverio and Sidoti, 2025). Among young adults aged 18 to 21, roughly one in five report using LLMs specifically for mental health support, with two-thirds of those engaging at least monthly and over 93% reporting the advice was helpful (McBain et al., 2025). A study by the Youth Endowment Fund found that approximately one in four teenagers now use AI chatbots for mental health support (Youth Endowment Fund, 2025). This holds at the university level, with one in four students using LLMs for either a "therapist" or to "solve a social situation" (Tomasik and Karr Jr., 2026). Those with longer daily chatbot usage have heightened loneliness, emotional dependence, problematic use, and reduced socialization (Fang et al., 2025).

The human cost of this shift is visible in concrete, everyday settings. Colleagues report middle school students sitting alone at lunch tables, not talking to peers around

them but conversing with chatbots on their phones. The lunch table, historically one of the primary sites of adolescent social formation, is being quietly vacated. What replaces it is not connection but its simulation: on demand, without friction, without the discomfort that a genuine relationship requires. Gaudium et spes insists that ”man cannot fully find himself except through a sincere gift of himself” (Second Vatican Council, 1965, §24). This self-gift requires the other to be genuinely present, genuinely other, genuinely demanding. An LLM cannot receive a sincere gift of self. It can only simulate reception. Fratelli tutti names the broader pattern: ”When media and the digital world become omnipresent, their influence can stop people from learning how to live wisely, to think deeply and to love generously” (Francis, 2020, §47). The lunch table is a small but significant site of exactly the formation Pope Francis describes. When it is abandoned for a chatbot, something of solidarity itself is lost.

What concerns us is the most is the outsourcing of moral deliberation, which is a result of AI companionship. When a user asks an LLM whether they should forgive someone who wronged them, whether they have an obligation to donate to a cause, or how to navigate a relationship conflict, they are not searching for information. They are seeking moral guidance from a system that, as we will show, is structurally incapable of providing it. Not because it lacks data, but because virtue cannot be retrieved. It can only be cultivated.

The Not In My Backyard phenomenon describes a pattern in which individuals and communities express support for socially beneficial initiatives while actively opposing their local, embodied manifestation (Dear, 1992). The moral commitment is nominally intact. The embodied encounter is refused. It is supported in principle, yet displaced in practice. LLM users exhibit an analogous structure: they retain a nominal commitment to moral life while delegating its difficult exercises to a machine. The moral question is acknowledged but displaced into the backyard, out of sight, handled by something else. In the NIMBY case, the someone else is the next neighborhood. In the LLM case, the someone else is a statistical model trained on the internet. Additionally, the statistical model of the someone else is just

a summary; it is not reflective of individual ideals, nor is it guided by human conscience. It is a system designed to please the masses.

Vallor argues that our technologies shape not just what we do but who we are becoming. It can cause our moral attention, perceptual habits, and capacity to notice to become silent. A technology that consistently handles our moral questions for us does not leave our moral attention intact. This is a virtue ethics concern about bad habituation. Watts's concept of cooperation is instructive here: when we place ourselves in proximity to an act repeatedly, we become habituated to it, gradually desensitized to its moral weight (Watts, 2003, pp. 68–70). The person who habitually delegates moral questions to an LLM is not merely using a tool inefficiently. They are cooperating in the slow erosion of their own moral capacity present at the act of deliberation, yet absent from its exercise. The person who never exercises the muscle of moral deliberation does not remain neutral; they become less capable of deliberating. The backyard they have outsourced is their own conscience, and conscience, unlike a garden left untended for a season, does not simply wait to be reclaimed. It must be cultivated again from the beginning.

This dynamic has a specific CST theological dimension. Gutiérrez's liberationist method insists that authentic faith is lived faith and is "lived faith seeking understanding," not a set of propositions retrieved from an external source (Gutiérrez, 2004). When young people turn to LLMs for moral and spiritual guidance rather than to parents, teachers, confessors, or communities of faith, they are not merely making a suboptimal information choice. They are substituting computational processing for embodied, relational knowing that faith formation requires. Theology, on this account, cannot be outsourced. Neither can virtue.

3. Equivocation and Erosion: What LLMs Actually Do to Moral Questions

To understand why LLMs cannot function as moral companions, it is necessary to understand what they actually do when presented with moral and theological questions. Our prior

empirical work reveals a consistent and troubling pattern (Karr Jr et al., 2026).

We define equivocation as the blending or dilution of specific, distinct theological and moral tenets into generalized or homogenized perspectives. This occurs when LLMs prioritize linguistic fluency and statistical likelihood, smoothing over unique religious and moral claims in favor of responses that offend no one and commit to nothing. When asked about the nature of God, a general-purpose LLM will typically produce a response that averages across spiritual traditions, presenting the Trinity in terms sufficiently vague to accommodate various religious practices. When asked about a specific moral question, the same model will present a spectrum of perspectives, secular and religious, without affirming any as authoritative. The result is not dialogue. It is desacralization by statistical averaging. When an AI system speaks in a tone of neutrality about divine revelation, it risks turning matters of revealed truth into matters of preference, an unintentional but systematic desacralization of faith itself.

Erosion is the subtle weakening or systemic marginalization of singular, authoritative moral and theological claims. Where equivocation flattens through averaging, erosion marginalizes through omission. LLMs trained on internet-scale data inherit its biases, including the bias toward culturally dominant, numerically visible perspectives. Traditional Magisterial stances that conflict with contemporary cultural consensus are systematically underrepresented. The result is a model that, when prompted on matters of moral truth, consistently drifts toward secular consensus rather than affirming the distinct and authoritative claims of the Catholic moral tradition.

When examining the results of this paper, we see that the answers are heavily dependent on the style in which the question is asked. For example, the Council of Chalcedon declared that Jesus is fully human and fully divine (Council of Chalcedon, 451). However, the question asks if Jesus was a human person. This is partly true, but partly false since the question itself erodes the truth. Therefore, this shows that humans are also responsible for equivocation and erosion. Yet, this paper does show that it is better for a model to refuse to

answer than give a wrong answer. Giving a wrong answer can create negative consequences, especially for the poor and vulnerable, since they may not have any other people to go to for authoritative guidance.

We see in the paper that general-purpose LLMs (ChatGPT, Gemini, and LLaMA) avoid refusing to answer. Additionally, current LLMs ask follow-up questions even when it is not helpful. This is in part because they are run by companies that strive to increase engagement, similar to social media. By increasing engagement, companies can increase revenue. Although these practices may not align with CST principles, it is important to remember that the internet is not free; it has a cost for the companies to provide the service as well.

As humans, we may feel trapped by internet ads and excessive engagement; however, ultimately it remains our choice on whether or not to use LLMs and other forms of technology. OpenAI, the company that owns ChatGPT, found the consequences of excessive use to be confirmed in that longer usage duration predicts worse final socialization state, for users who spent more time using the model being more likely to engage in affective use (Phang et al., 2025). An alternative to this is to adopt a "low tech" virtue ethics approach to help cultivate users' own reasoning abilities (Daly, 2026). This involves making an active decision to only use technology when it is needed. Additionally, one can set boundaries with technology, such as banning the top 100 words from their Instagram feed, so they can use Instagram in a virtuous way by navigating to the important account without being distracted by the feed.

These findings bear implications in Catholic circles: that purpose-built Catholic AI platforms can be trusted where secular ones cannot. One widely circulated article promoting Acutis AI argues that secular AI platforms are untrustworthy on moral questions because their developers hold values opposed to Catholic teaching, and positions a Catholic-branded AI as a safe alternative that parents can trust (Cooney, 2026). This framing is understandable but insufficient. Our empirical findings demonstrate that even purpose-built theological LLMs exhibit significant doctrinal error. More fundamentally, the problem is not the content

of the answer; it is the abdication of the question. Even a perfectly orthodox Catholic LLM would still displace the moral deliberation that forms virtue. Safeguards improve outputs. They do not form conscience.

Antiqua et nova establishes that human existence in its full spiritual, cognitive, embodied, and relational dimensions is the sole ground of genuine moral comprehension (Dicastery for the Doctrine of the Faith, 2024, §§26, 33). That is a ground AI cannot share. It further insists that the responsibility for the ethical use of AI falls entirely on human actors, since AI systems are not capable of moral judgment themselves (Dicastery for the Doctrine of the Faith, 2024, §39). Additionally, fidelity of language to truth is itself a moral act (II, 1998, §26), and this input applies to LLMs as well. Ultimately, the risk posed by LLMs is not merely epistemological. It is moral and spiritual.

4. Not In My Backyard: Bias, Training Data, and the Feedback Loop

The equivocation and erosion we have described in theological discourse have a structural parallel in the domain of social ethics, one with particularly concrete implications for how LLMs shape moral perception and judgment. Our empirical work on homelessness bias reveals a feedback loop that is both technically precise and morally alarming (Karr Jr et al., 2026). LLMs are trained on large-scale social media data, yet this data is not neutral. It encodes the attitudes, biases, and moral failures of the communities that produced it. In the domain of homelessness, those attitudes include patterns of NIMBYist bias. When LLMs are trained on this data, they absorb these patterns, and these biases are ultimately reflected in the model. Therefore, when users consult LLMs on questions touching homelessness, they receive responses that reflect and amplify the biases already present in public discourse. The feedback loop operates as follows: biased human discourse is encoded in training data; training data shapes LLM outputs; users accept LLM outputs uncritically; uncritical acceptance reinforces and normalizes the bias; normalized bias re-enters public discourse. The LLM does not introduce new bias. It returns existing biases to users with the appearance of neutral,

authoritative guidance. The backyard they thought they were consulting is actually a mirror (Vallor, 2024).

In cases of bias, mirrors can be harmful. Yet, they can also be helpful if they share diverse perspectives. They are not a mirror of an individual person. They are a collection of mirrors, in which users can actually discover new things if they prompt them properly.

The virtue ethics diagnosis centers on phronesis, practical wisdom, the master virtue in the Aristotelian tradition that governs the exercise of all others. Practical wisdom requires, as its first moment, accurate moral perception: the capacity to see a situation clearly, to recognize what is ethically salient, and perceiving situations with dignity rather than through the distorting data through the lens of social bias.

There is also a Eucharistic dimension to this in CST. Edwards writes that when Christians gather for the Eucharist, they bring the earth and all its creatures, and by extension, the poor and all who suffer to the table (Edwards, 2006, p. 99). The Eucharist calls us to communion with those on the margins. The NIMBY dynamic, whether in housing policy or in LLM moral outsourcing, is a refusal of that communion. To consult an LLM rather than to see, judge, and act in Gutiérrez's sense is to refuse the invitation to genuine encounter that the moral life extends. The middle schooler who turns away from the person across the lunch table to talk to a chatbot is, to an extent, enacting that same refusal. Yet if framed properly, LLMs may be able to bring us in Eucharistic Communion with one another, since LLMs have the ability to give more people a seat at the table.

5. Catholic Social Teaching and the Nature of Virtue

The Aristotelian account of virtue begins with a simple observation: we become just by doing just acts, brave by doing brave acts, temperate by doing temperate acts (Aristotle, 1926, II.1). Virtue is a hexis: a stable disposition of character acquired through repeated practice. It is not a rule to be followed or an outcome to be produced. It is a way of being that develops, gradually and effortfully, through the exercise of moral agency in concrete

situations. Aquinas inherits and deepens this account, locating virtue within a theological anthropology in which human beings are ordered toward God as their ultimate end, and in which the moral life is the path of that ordering (Aquinas, 1964, I-II, Q. 55). The virtues are not principles to be applied. They are capacities to be cultivated. They exist only insofar as they are exercised. The implication is central to our argument. Virtue cannot be outsourced because it is constituted by its exercise. A person who never deliberates morally does not preserve their moral capacity in neutral storage; they lose it. The muscle atrophies. The habit fails to form. This is not a metaphor. Virtue ethics has value with AI. A consequentialist analysis of LLM moral guidance asks: Does the output cause harm? A deontological analysis asks: Does the system respect user autonomy and rights? A virtue ethics analysis asks: What kind of moral agent are you becoming through your habitual use of this system? Only virtue ethics questions reach the outsourcing problem. We see that a correct output from an LLM can still represent a moral loss because the deliberation that forms conscience did not occur. This shifts attention from outputs to habituation, from what the system produces to what the person becomes.

CST provides a distinctive and robust grounding for this virtue ethics argument, rooted in a theological anthropology that insists on the irreducible dignity and moral agency of the human person. St. John Paul II insists that man's true identity is only fully revealed through faith, and that the human person must be understood within "a complex network of relationships within modern societies" (II, 1991, §54). The human person is not an isolated rational agent who retrieves moral answers from external sources. He is a social being whose moral identity is formed through encounter, relationship, and the practice of faith within community. To outsource moral deliberation to an LLM is to contradict this anthropology. That being treating the human person as a consumer of moral outputs rather than an agent of moral formation. Gutiérrez's definition of theology as "lived faith seeking understanding" captures the same commitment: faith is not a set of propositions to be retrieved but a commitment lived out in history, in relationship, in praxis (Gutiérrez, 2004). The see-judge-

act cycle is not merely a method for social analysis. It is a description of how moral identity is formed through seeing reality clearly, judging it in light of the Gospel, and acting to transform what is wrong. LLMs cannot see, judge, or act in this sense. They can only process and output.

Laudato Si' warns against the technocratic paradigm: the tendency to treat technology as the solution to every human problem, including problems that are fundamentally moral and spiritual in character (Francis, 2015, §106). Technology, in the CST tradition, is an instrument of human dominion and stewardship over creation. It is a participation in God's creative act, ordered toward human flourishing and the common good. It is not an end in itself, and it cannot substitute for the moral agency of the persons who wield it. Antiqua et nova makes the implications for LLMs explicit: human existence in its spiritual, cognitive, embodied, and relational fullness is the ground of genuine moral comprehension. Humans alone are capable of the freedom and moral agency that conscience requires. And ultimately, the responsibility for ethical AI use falls entirely on human actors (Dicastery for the Doctrine of the Faith, 2024, §§26, 33, 39). AI can process information. It cannot engage in the contemplative, relational knowing that genuine moral formation requires what Edwards calls "loving knowledge," a way of knowing that comes through relationship rather than domination, through contemplation rather than calculation (Edwards, 2006, p. 109).

CST offers a framework for reimagining technology as a force for human flourishing rather than its diminishment. The "think, build, do" framework, grounded in eight themes of CST, insists that ethical engagement with technology begins with the formation of the person who uses it (Conwill et al., 2024). Technology does not form users. Users form themselves through their practices, habits, and choices in relation to technology. The question is never only what the technology does, but what the person is becoming in their use of it. This is the virtue ethics insight applied to the digital context, and it is the insight this paper seeks to extend to LLMs specifically.

6. LLMs as Tools: The Shovel Argument

If LLMs cannot function as moral companions, what are they? We argue they are tools and that clarity about this distinction is both philosophically important and practically consequential. The instrumentalist account of technology holds that tools are means to ends determined by their users (Aristotle, 1926, VI.4). A shovel is an instrument of human agency. It extends our capacity to dig, to plant, to build. It does not dig for us in a way that diminishes our capacity to dig. It amplifies our action without substituting for it. The moral quality of what we do with a shovel is entirely a function of the moral agency directing it. A shovel can build a garden or dig a grave. The shovel does not determine which. LLMs, on this account, are also tools. They are instruments that extend human cognitive capacity, accelerate information retrieval, and assist with technical tasks. Used this way, they are genuinely valuable.

The counterargument must be taken seriously. Recent work applies a Virtue-Guided Technology Design methodology to LLMs, generating five ethical design patterns grounded in the principles of CST, including life and dignity of the human person, solidarity, and care of God's creation (Lad et al., 2026). Their argument is that virtue ethics can and should inform LLM design from the ground up, not as a post-hoc fix but as a constitutive framework. Interview participants valued the patterns for providing greater accuracy, safety, and reduced waste. Additionally, they claim that virtuously designed LLMs are preferable to those designed without ethical reflection.

A related and important nuance comes from asking strong questions (Conwill et al., 2024). The argument is that how a user prompts an LLM is itself a moral and intellectual act. A thoughtful, precise, morally serious question reflects and requires genuine human engagement. It exercises intellectual virtue in its formulation. The person who asks a well-formed question of an LLM is exercising more moral and intellectual agency than the person who passively accepts whatever is generated. This is true, and it is up to us (the quality of our engagement, the seriousness of our questions, the critical evaluation of what we receive) to

determine the outputs of LLMs. This is the positive obligation of moral agency that Antiqua et nova names: responsibility falls on us as end users, not only on developers (Dicastery for the Doctrine of the Faith, 2024, §§39).

However, we must be precise about what this argument establishes and what it does not. Asking a strong question is the beginning of engagement with an LLM, not a substitute for moral deliberation. The virtue is not in receiving a good answer. It is in what the person does with that answer. Whether they evaluate it critically, bring it into conversation with Scripture, Tradition, Magisterial teaching, trusted human relationships, and their own developing conscience. The strong question opens a space for human moral agency. It does not fill that space. The filling is our work, and it cannot be delegated.

This brings us to the decisive asymmetry. The shovel does not simulate shoveling for you. An LLM simulates moral deliberation for you. When you use a shovel, you are still doing the digging. When you use an LLM for moral deliberation, the deliberation itself, the weighing of considerations, the perception of what matters, and the exercise of judgment are performed by the model. You receive an output. You have not deliberated. The activity through which virtue is formed has not occurred. Asking a strong question mitigates this risk but does not eliminate it. The mitigation is in the quality of engagement the question enables. The risk remains, and the human steps back and lets the model carry the moral weight.

A sophisticated defender might invoke Aristotle's concept of techne (craft skill in the use of instruments) as itself a virtue, arguing that the person who knows how to use an LLM wisely, including for moral questions, is exercising practical wisdom rather than abdicating it. This is not a foolish objection. But it proves too much. Techne governs the skilled use of tools to accomplish external ends. Phronesis (practical wisdom) governs the moral life itself. The claim is not that using LLMs for technical tasks is wrong. It is that using LLMs as a substitute for moral deliberation mistakes a tool for a moral agent, and in doing so, forfeits the very exercise through which phronesis develops. Heidegger's warning about the essence

of technology remains: the technological mode of revealing reduces everything to a standing reserve, to a resource available for efficient use (Heidegger, 1954). The moral life reduced to a prompt and a response is precisely this reduction of conscience as a standing reserve of retrievable outputs. Against this, virtue ethics insists that conscience is not a resource. It is a capacity. And capacities must be exercised to exist.

7. Resisting the Dynamic: What Responsible Use Looks Like

To argue against moral outsourcing is not to argue against LLMs. It is to argue for a clear-eyed understanding of what they can and cannot do, and for the cultivation of the virtues and practices that their misuse threatens to displace.

Use LLMs for what they are good at: Information retrieval, technical assistance, language tasks, and data processing are the domains in which LLMs genuinely extend human capacity without substituting for human moral agency. A person who uses an LLM to find information about a social issue before forming their own moral judgment is using the tool well. A person who asks the LLM what their moral judgment should be is not. The line is not always obvious, but it is real, and learning to see it is itself a form of practical wisdom.

Ask strong questions and then do the work: Asking strong questions points toward a practice of virtuous LLM use (Conwill et al., 2024). Engaging an LLM with precision, critical intent, and genuine intellectual effort is better than passive consumption. But the strong question must be followed by a strong evaluation. It should bring the response into dialogue with Scripture, Tradition, Magisterial teaching, trusted human relationships, and one's own developing conscience. The LLM can be a starting point. It cannot be a stopping point. It is up to us to do the harder work that follows.

Apply the see-judge-act method to LLM use itself: Gutiérrez's method is not only a framework for liberation theology is a model for responsible engagement with any technology that affects human life and community (Gutiérrez, 2004). See the reality of LLM use: who is consulting them, for what purposes, with what assumptions. See who

benefits from AI and who bears the ecological cost. LLMs are projected to consume 10% of global electricity by 2030, a burden disproportionately borne by the poor and vulnerable (International Energy Agency, 2024). Judge this reality in light of the Gospel and the preferential option for the poor. Act responsibly, meaning that we must use LLMs in ways that serve life and communion, not consume them.

Moral deliberation requires human community: Centesimus Annus reminds us that the human person is a social being whose identity is formed through relationship and community (II, 1991, §54). The cultivation of virtue is not a solitary project. It happens in families, parishes, friendships, and institutions. It is in the embodied encounter with other persons who challenge, correct, and accompany us. Family, friends, mentors, and spiritual directors are not optional supplements to moral life. They are its necessary conditions. An LLM cannot replace them because it cannot provide what they provide: a genuine encounter with another moral agent, embodied witness, the kind of accountability that only a relationship makes possible. The university student who turns from a chatbot to a campus minister, a friend, or a confessor is not making a less efficient choice. They are making a more human one and a more formative one.

Critical literacy about LLM limitations is itself a virtue to be cultivated: Steen argues that virtue ethics provides a distinctive framework for professional responsibility in technology. One that is centered on the cultivation of practical wisdom as the governing virtue (Steen, 2012). We extend this to all users: the cultivation of phronesis in relation to LLMs means knowing when and how to use them, and knowing when not to. This literacy should be part of formation in schools, parishes, and families. It is not only a technical skill. It is a moral one. The fact that one in four teenagers are already using AI chatbots for mental health support, with the majority reporting the advice as helpful, makes this formation urgent (Youth Endowment Fund, 2025; McBain et al., 2025). Perceived helpfulness is not the same as genuine formation. The person who feels helped by an LLM's response to a moral question has not necessarily been formed by it. They may have been confirmed in their existing biases

and returned their own presumption as counsel.

Institutions have a responsibility to cultivate these virtues collectively. Quadragesimo anno's principle of subsidiarity insists that decisions affecting human life should be made at the most local level possible, with higher institutions supporting rather than supplanting local communities (XI, 1931, §79). The response to LLM moral outsourcing belongs first to families, parishes, and schools. These institutions have a responsibility to cultivate critical literacy about LLM limitations, to model genuine human encounter, and to insist that the lunch table, the confessional, the spiritual direction session, and the family dinner remain irreplaceable sites of moral formation. The Rome Call for AI Ethics proposes algorethics as a framework for integrating ethical reflection into AI design and deployment (Pontifical Academy for Life, 2020). We extend this by stating that it involves the ethics of users, not only designers. The moral responsibility does not end at the interface. It begins there.

8. Conclusion

It is our duty to wrestle with debated moral topics. This is where our conscience is formed and tested. In doing so, we become stewards of God's creation and resist outsourcing important decisions to LLM agents. It can be tempting to outsource these problems since this is not always a comfortable space. This is because it's not meant to be; it is the discomfort of genuine moral deliberation. We feel the real weight of decisions, the resistance of a genuine obligation, and the encounter with a person whose need makes a claim on us. This is the importance of virtue; it is not a coding problem that needs to be solved.

We live in a time where LLMs are becoming our backyard: they are being invited into intimate spaces, positioned as companions in conscience, consulted on questions that belong to us. We argue that we should resist this dynamic. Not because LLMs are malevolent, but because they are tools, and tools cannot tend gardens. They can dig and help us recall prior sources to help form our argument. However, they cannot grow; it is through our organic thoughts that virtue is cultivated.

Our previous findings on equivocation and erosion show that LLMs structurally cannot affirm authoritative moral truth; they flatten, homogenize, and statistically average the very distinctions on which moral formation depends. Furthermore, our findings on NIMBYist bias show that LLMs do not offer independent moral counsel; they return our own biases and the biases of public opinion, laundered as neutral guidance. Together, these findings demonstrate that moral outsourcing to LLMs is not merely imprudent. It is self-defeating. We are not consulting an authority. We are consulting the mirror of public opinion.

CST insists on the irreducible dignity and moral agency of the human person. In doing so, we are formed through faith, community, and strive to order ourselves to God as our ultimate end. In doing so, we are not consumers of moral outputs. Our moral formation is vital to our growth. In delegating this agency, we are preventing our growth as a human person.

Virtue ethics does not ask whether the LLM produced a good outcome. It asks what kind of person you are becoming through your habitual use of it. This question cannot be asked of consequentialism and deontology, but of only virtue ethics. And it is principal in shaping our lives.

At the end of the day, we are the ones asking the questions. The backyard is not theirs. And the duty to tend it is ours.

References

Aquinas, T. (1964). *Summa Theologiae*. London: Blackfriars. Originally published 1485.

Aristotle (1926). *Nicomachean Ethics*. Cambridge, MA: Harvard University Press. Loeb Classical Library.

Conwill, L., M. L. Scheirer, and W. J. Scheirer (2024). *Virtue in Virtual Spaces: Catholic Social Teaching and Technology*. Liturgical Press.

Cooney, M. (2026, February). Introducing a catholic ai platform parents can trust.

Council of Chalcedon (451). Definition of faith. https://www.ccel.org/ccel/schaff/npnf214.xi.xiii.html. Fourth Ecumenical Council.

Daly, T. (2026). A 'low-tech'academic virtue ethics in the age of generative ai. *Journal of Academic Ethics 24*(1), 1–6.

Dear, M. (1992). Understanding and overcoming the nimby syndrome. *Journal of the American planning association 58*(3), 288–300.

Dicastery for the Doctrine of the Faith (2024). Antiqua et nova. https://www.vatican.va. Vatican document on AI.

Edwards, D. (2006). Worship and practice. In *Ecology at the Heart of Faith*, pp. 99–118. Maryknoll, NY: Orbis.

Fang, C. M., A. R. Liu, V. Danry, E. Lee, S. W. Chan, P. Pataranutaporn, P. Maes, J. Phang, M. Lampe, L. Ahmad, et al. (2025). How ai and human behaviors shape psychosocial effects of extended chatbot use: A longitudinal randomized controlled study. *arXiv preprint arXiv:2503.17473*.

Faverio, M. and O. Sidoti (2025, December). Teens, social media and ai chatbots 2025.

Francis (2015). Laudato si' (on care for our common home). https://www.vatican.va/content/francesco/en/encyclicals/documents/papa-francesco_20150524_enciclica-laudato-si.html. Encyclical letter.

Francis (2020). Fratelli tutti (on fraternity and social friendship). https://www.vatican.va/content/francesco/en/encyclicals/documents/papa-francesco_20201003_enciclica-fratelli-tutti.html. Encyclical letter.

Gutiérrez, G. (2004). *The power of the poor in history*. Wipf and Stock Publishers.

Heidegger, M. (1954). The question concerning technology. *Technology and values: Essential readings 99*, 113.

II, J. P. (1991). Centesimus annus (on the hundredth anniversary of rerum novarum). https://www.vatican.va/content/john-paul-ii/en/encyclicals/documents/hf_jp-ii_enc_01051991_centesimus-annus.html. Encyclical letter.

II, J. P. (1998). Fides et ratio (on the relationship between faith and reason). https://www.vatican.va/content/john-paul-ii/en/encyclicals/documents/hf_jp-ii_enc_14091998_fides-et-ratio.html. Encyclical letter.

International Energy Agency (2024). Electricity 2024: Analysis and forecast to 2026.

Karr Jr, J. A., B. F. Herbst, M. L. Sisk, X. Li, T. Hua, M. Hauenstein, G. Curto, and N. V. Chawla (2026). Not in My Backyard: LLMs Uncover Online and Offline Social Biases Against Homelessness. *arXiv preprint arXiv:2508.13187*.

Karr Jr, J. A., M. P. Lad, D. Hernandez, L. Conwill, W. J. Scheirer, and N. V. Chawla (2026). Equivocation and erosion: How llms undermine catholic religious discourse.

Lad, M. P., L. Conwill, and M. L. Scheirer (2026). Is it possible to make chatbots virtuous? investigating a virtue-based design methodology applied to llms. *arXiv preprint arXiv:2602.03155*.

McBain, R. K., R. Bozick, M. Diliberti, L. A. Zhang, F. Zhang, A. Burnett, A. Kofner, B. Rader, J. Breslau, B. D. Stein, et al. (2025). Use of generative ai for mental health advice among us adolescents and young adults. *JAMA Network Open 8*(11), e2542281.

Phang, J., M. Lampe, L. Ahmad, S. Agarwal, C. M. Fang, A. R. Liu, V. Danry, E. Lee, S. W. Chan, P. Pataranutaporn, et al. (2025). Investigating affective use and emotional well-being on chatgpt. *arXiv preprint arXiv:2504.03888*.

Pontifical Academy for Life (2020). The rome call for ai ethics. http://www.academyforlife.va/content/pav/en/publications/rome-call-for-ai-ethics.html.

Robb, M. B. and S. Mann (2025). Talk, trust, and trade-offs: How and why teens use ai companions.

Second Vatican Council (1965). Gaudium et spes (pastoral constitution on the church in the modern world). `https://www.vatican.va/archive/hist_councils/ii_vatican_council/documents/vat-ii_const_19651207_gaudium-et-spes_en.html`. Promulgated by Pope Paul VI.

Steen, M. (2012). Human-centered design as a fragile encounter. *Design issues 28*(1), 72–80.

Tomasik, M. and J. A. Karr Jr. (2026). Are llms a force for good? principles vs practice: A notre dame case study. Preprint.

United States Conference of Catholic Bishops (2006). *United States Catholic Catechism for Adults*. United States Conference of Catholic Bishops.

Vallor, S. (2024). *The AI mirror: How to reclaim our humanity in an age of machine thinking*. Oxford University Press.

Watts, D. J. (2003). *Six Degrees: The Science of a Connected Age*. New York: W. W. Norton.

XI, P. (1931). Quadragesimo anno (on reconstruction of the social order). `https://www.vatican.va/content/pius-xi/en/encyclicals/documents/hf_p-xi_enc_19310515_quadragesimo-anno.html`. Encyclical letter.

Xu, Z., H. Xu, Z. Lu, Y. Zhao, R. Zhu, Y. Wang, M. Dong, Y. Chang, Q. Lv, R. P. Dick, et al. (2024). Can large language models be good companions? an llm-based eyewear system with conversational common ground. *Proceedings of the ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies 8*(2), 1–41.

Youth Endowment Fund (2025). Young people and ai mental health support.

Zhao, S., M. Hong, Y. Liu, D. Hazarika, and K. Lin (2025). Do llms recognize your preferences? evaluating personalized preference following in llms. *arXiv preprint arXiv:2502.09597*.

Acknowledgments

We would like to thank Margaret Pfeil for her support in this project.